

The purpose of this project is to write a program to compute the optimal sequence alignment of two DNA strings. This program will introduce you to the emerging field of *computational biology* in which computers are used to do research on biological systems. Further, you will be introduced to a powerful algorithmic design paradigm known as *dynamic programming*.

Biology Review A *genetic sequence* is a string formed from a four-letter alphabet Adenine (A), Thymine (T), Guanine (G), Cytosine (C) of biological macromolecules referred to together as the DNA bases. A *gene* is a genetic sequence that contains the information needed to construct a protein. All of your genes taken together are referred to as the human genome, a blueprint for the parts needed to construct the proteins that form your cells. Each new cell produced by your body receives a copy of the genome. This copying process, as well as natural wear and tear, introduces a small number of changes into the sequences of many genes. Among the most common changes are the substitution of one base for another and the deletion of a substring of bases; such changes are generally referred to as *point mutations*. As a result of these point mutations, the same gene sequenced from closely related organisms will have slight differences.

The Problem Through your research you have found the following sequence of a gene in a previously unstudied organism.

A A C A G T T A C C

What is the function of the protein that this gene encodes? You could begin a series of uninformed experiments in the lab to determine what role this gene plays. However, there is a good chance that it is a variant of a known gene in a previously studied organism. Since biologists and computer scientists have laboriously determined (and published) the genetic sequence of many organisms (including humans), you would like to leverage this information to your advantage. We'll compare the above genetic sequence with one which has already been sequenced and whose function is well understood.

T A A G G T C A

If the two genetic sequences are similar enough, we might expect them to have similar functions. We would like a way to quantify "similar enough."

Edit Distance In this assignment we will measure the similarity of two genetic sequences by their *edit distance*, a concept first introduced in the context of coding theory, but which is now widely used in spell checking, speech recognition, plagiarism detection, file revisioning, and computational linguistics. We align the two sequences, but we are permitted to *insert gaps* in either sequence (eg, to make them have the same length). We pay a penalty for each gap that we insert and also for each pair of characters that mismatch in the final alignment. Intuitively, these penalties model the relative likeliness of point mutations arising from deletion/insertion and substitution. We produce a numerical score according to the following table, which is widely used in biological applications:

operation	cost
insert a gap	2
align two characters that mismatch	1
align two characters that match	0

Here are two possible alignments of the strings $x = \text{'AACAGTTACC'}$ and $y = \text{'TAAGGTCA'}$:

x	y	cost	x	y	cost
---	---	---	---	---	---
A	T	1	A	T	1
A	A	0	A	A	0
C	A	1	C	-	2
A	G	1	A	A	0
G	G	0	G	G	0
T	T	0	T	G	1
T	C	1	T	T	0
A	A	0	A	-	2
C	-	2	C	C	0
C	-	2	C	A	1
---	---	---	---	---	---
		8			7

The first alignment has a score of 8, while the second one has a score of 7. The *edit distance* is the score of the best possible alignment between the two genetic sequences over all possible alignments. In this example, the second alignment is in fact optimal, so the edit distance between the two strings is 7. Computing the edit distance is a nontrivial computational problem because we must find the best alignment among exponentially many possibilities. For example, if both strings are 100 characters long, then there are more than 10^{75} possible alignments.

We will explain a recursive solution which is an elegant approach. However it is far too inefficient because it recalculates each subproblem over and over. Once we have defined the recursive definition we can redefine the solution using a dynamic programming approach which calculates each subproblem once.

A Recursive Solution We will calculate the edit distance between the two original strings x and y by solving many edit-distance problems on smaller suffixes of the two strings. We use the notation $x[i]$ to refer to character i of the string. We also use the notation $x[i..M]$ to refer to the suffix of x consisting of the characters $x[i]$, $x[i + 1]$, ..., $x[M - 1]$. Finally, we use the notation $\text{opt}[i][j]$ to denote the edit distance of $x[i..M]$ and $y[j..N]$. For example, consider the two strings $x = \text{'AACAGTTACC'}$ and $y = \text{'TAAGGTCA'}$ of length $M = 10$ and $N = 8$, respectively. Then, $x[2]$ is 'C', $x[2..M]$ is 'CAGTTACC', and $y[8..N]$ is the empty string. The edit distance of x and y is $\text{opt}[0][0]$.

Now we describe a recursive scheme for computing the edit distance of $x[i..M]$ and $y[j..N]$. Consider the first pair of characters in an optimal alignment of $x[i..M]$ with $y[j..N]$. There are three possibilities:

1. The optimal alignment matches $x[i]$ up with $y[j]$. In this case, we pay a penalty of either 0 or 1, depending on whether $x[i]$ equals $y[j]$, plus we still need to align $x[i + 1..M]$ with $y[j + 1..N]$. What is the best way to do this? This subproblem is exactly the same as the original sequence alignment problem, except that the two inputs are each suffixes of the original inputs. Using our notation, this quantity is $\text{opt}[i + 1][j + 1]$.
2. The optimal alignment matches the $x[i]$ up with a gap. In this case, we pay a penalty of 2 for a gap and still need to align $x[i + 1..M]$ with $y[j..N]$. This subproblem is identical to the original sequence alignment problem, except that the first input is a proper suffix of the original input.
3. The optimal alignment matches the $y[j]$ up with a gap. In this case, we pay a penalty of 2 for a gap and still need to align $x[i..M]$ with $y[j + 1..N]$. This subproblem is identical to the original sequence alignment problem, except that the second input is a proper suffix of the original input.

The key observation is that all of the resulting subproblems are sequence alignment problems on suffixes of the original inputs. To summarize, we can compute $\text{opt}[i][j]$ by taking the minimum of three quantities:

$$\text{opt}[i][j] = \min\{\text{opt}[i + 1][j + 1] + 0 \text{ or } 1, \text{opt}[i + 1][j] + 2, \text{opt}[i][j + 1] + 2\}$$

This equation works assuming $i < M$ and $j < N$. Aligning an empty string with another string of length k requires inserting k gaps, for a total cost of $2k$. Thus, in general we should set $\text{opt}[M][j] = 2(N - j)$ and $\text{opt}[i][N] = 2(M - i)$. For our example, the final matrix is:

		0	1	2	3	4	5	6	7	8
$x \backslash y$		T	A	A	G	G	T	C	A	-
0	A	7	8	10	12	13	15	16	18	20
1	A	6	6	8	10	11	13	14	16	18
2	C	6	5	6	8	9	11	12	14	16
3	A	7	5	4	6	7	9	11	12	14
4	G	9	7	5	4	5	7	9	10	12
5	T	8	8	6	4	4	5	7	8	10
6	T	9	8	7	5	3	3	5	6	8
7	A	11	9	7	6	4	2	3	4	6
8	C	13	11	9	7	5	3	1	3	4
9	C	14	12	10	8	6	4	2	1	2
10	-	16	14	12	10	8	6	4	2	0

By examining $\text{opt}[0][0]$, we conclude that the edit distance of x and y is 7.

Problem 1. (*Calculating Edit Distance Using Dynamic Programming*) A direct implementation of the above recursive scheme will work, but it is spectacularly inefficient. If both input strings have N characters, then the number of recursive calls will exceed 2^N . To overcome this performance bug, we use *dynamic programming*. Dynamic programming is a powerful algorithmic paradigm, first introduced by Bellman in the context of operations research, and then applied to the alignment of biological sequences by Needleman and Wunsch. Dynamic programming now plays the leading role in many computational problems, including control theory, financial engineering, and bioinformatics, including BLAST (the sequence alignment program almost universally used by molecular biologists in their experimental work). The key idea of dynamic programming is to break up a large computational problem into smaller subproblems, *store* the answers to those smaller subproblems, and, eventually, use the stored answers to solve the original problem. This avoids recomputing the same quantity over and over again. Instead of using recursion, use a nested loop that calculates $\text{opt}[i][j]$ in the right order so that $\text{opt}[i + 1][j + 1]$, $\text{opt}[i + 1][j]$, and $\text{opt}[i][j + 1]$ are all computed *before* we try to compute $\text{opt}[i][j]$.

Write a program `edit_distance.py` that reads strings x and y from standard input and computes the edit-distance matrix opt as described above. The program should output x , y , the dimensions (number of rows and columns) of opt , and opt itself, using the following format:

- The first and second lines should contain the strings x and y .
- The third line should contain the dimensions of the opt matrix, separated by a space.
- The following lines should contain the rows of the opt matrix, each ending in a newline character. Use `stdio.write()` with the format string `'%3d'` to write out the elements of the matrix.

```
$ python3 edit_distance.py < data/example10.txt
AACAGTTACC
TAAGGTCA
11 9
 7  8 10 12 13 15 16 18 20
 6  6  8 10 11 13 14 16 18
 6  5  6  8  9 11 12 14 16
 7  5  4  6  7  9 11 12 14
 9  7  5  4  5  7  9 10 12
 8  8  6  4  4  5  7  8 10
 9  8  7  5  3  3  5  6  8
11  9  7  6  4  2  3  4  6
13 11  9  7  5  3  1  3  4
14 12 10  8  6  4  2  1  2
16 14 12 10  8  6  4  2  0
```

Problem 2. (*Recovering the Alignment*) Now that we know how to compute the edit distance between two strings, we next want to recover the optimal alignment itself. The key idea is to retrace the steps of the dynamic programming algorithm backwards, re-discovering the path of choices (highlighted in red in the table above) from $\text{opt}[0][0]$ to $\text{opt}[M][N]$. To determine the choice that led to $\text{opt}[i][j]$, we consider the three possibilities:

1. The optimal alignment matches $x[i]$ up with a gap. In this case, we must have $\text{opt}[i][j] = \text{opt}[i + 1][j] + 2$.
2. The optimal alignment matches $y[j]$ up with a gap. In this case, we must have $\text{opt}[i][j] = \text{opt}[i][j + 1] + 2$.
3. The optimal alignment matches $x[i]$ up with $y[j]$. In this case, we must have $\text{opt}[i][j] = \text{opt}[i + 1][j + 1]$ if $x[i]$ equals $y[j]$, or $\text{opt}[i][j] = \text{opt}[i + 1][j + 1] + 1$ otherwise.

Write a program `alignment.py` that reads from standard input, the output produced by `edit_distance.py`, ie, input strings x and y , and the opt matrix. The program should then recover an optimal alignment using the procedure described above, and write to standard output the edit distance between x and y and the alignment itself, using the following format:

- The first line should contain the edit distance, preceded by the text `'Edit distance = '`.
- Each subsequent line should contain a character from the first string, followed by the paired character from the second string, followed by the associated penalty. Use the character `'-'` to indicate a gap in either string.

```
$ python3 edit_distance.py < data/example10.txt | python3 alignment.py
Edit distance = 7
A T 1
A A 0
C - 2
A A 0
G G 0
T G 1
T T 0
A - 2
C C 0
C A 1
```

Data The data directory contains short test data files and actual genomic data files. Be sure to test your programs thoroughly using the short test files and the longer actual data files. Here are the optimal edit distances of several of the supplied files:

```
ecoli2500.txt    118
ecoli5000.txt    160
fli8.txt         6
fli9.txt         4
fli10.txt        2
ftsa1272.txt     758
gene57.txt       8
stx1230.txt      521
stx19.txt        10
stx26.txt        17
stx27.txt        19
```

and here are the test cases with unique optimal alignments:

```
$ python3 edit_distance.py < data/endgaps7.txt | python3 alignment.py
Edit distance = 4
a - 2
t t 0
a a 0
t t 0
t t 0
a a 0
t t 0
- a 2
```

```
$ python3 edit_distance.py < data/fli10.txt | python3 alignment.py
Edit distance = 2
T T 0
G G 0
G G 0
C T 1
G G 0
G G 0
A T 1
A A 0
C C 0
T T 0
```

Files to Submit

1. edit_distance.py
2. alignment.py
3. report.txt

Before you submit:

- Make sure your programs meet the input and output specifications by running the following command on the terminal:

```
$ python3 run_tests.py -v [<problems>]
```

where the optional argument `<problems>` lists the problems (`Problem1`, `Problem2`, etc.) you want to test, separated by spaces; all the problems are tested if no argument is given.

- Make sure your programs meet the style requirements by running the following command on the terminal:

```
$ pycodestyle <program>
```

where `<program>` is the `.py` file whose style you want to check.

- Make sure your report isn't too verbose, doesn't contain lines that exceed 80 characters, and doesn't contain spelling/grammatical mistakes.

Acknowledgements This project is an adaptation of the Global Sequence Alignment assignment developed at Princeton University by Thomas Clarke, Robert Sedgewick, Scott Vafai, and Kevin Wayne.